



BBRT molecular optimization: trained Transformer vs. Claude (Sonnet 4.6, Opus 4.8+thinking)

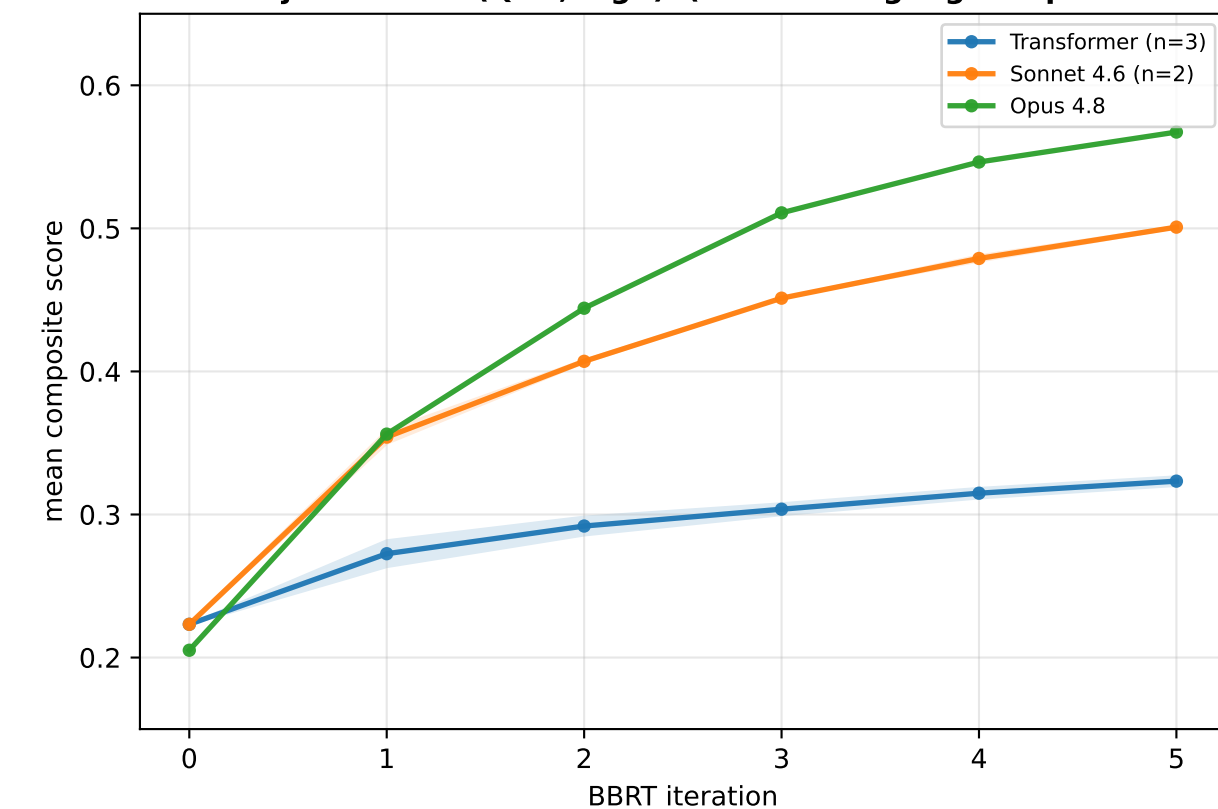
A. Penalized logP (unbounded → gameable)



B. QED (bounded [0,1] → honest)



C. Multi-objective: min(QED, logP) (natural-language 'improve BOTH')



D. Best molecule (△ hatched = reward-hacking artifact)

